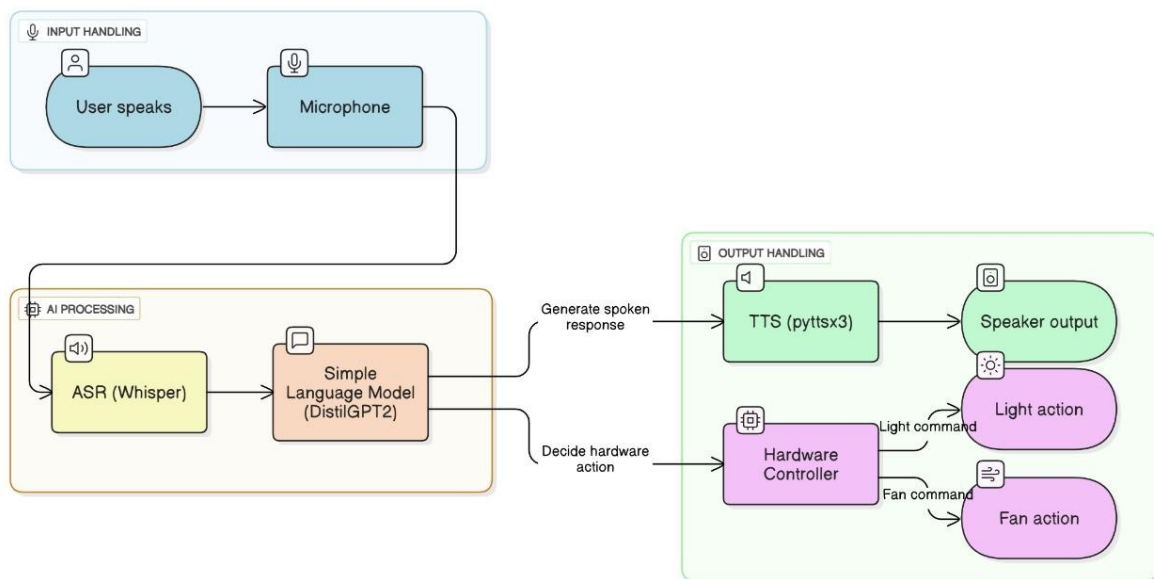


Research & Design Proposal

Local AI Voice Assistant

1) System Architecture



Explanation: -

- **User Voice:** The user speaks commands like “Turn on the light.”
- **ASR (Whisper):** Converts recorded speech to text.
- **SLM (DistilGPT2):** Processes text command, generates response text, and decides the hardware action.
- **TTS (pyttsx3):** Converts response text to audio for playback.
- **Speaker Output:** Plays the response audio.
- **Hardware Controller:** Executes the action (light/fan) on GPIO pins; in your implementation, this is simulated.

2) Model & Library Choices

| Stage | Model / Library | Notes / Rationale |
|----------|----------------------------|--|
| ASR | Whisper (small) | Accurate, offline transcription, lightweight enough for local CPU. |
| SLM | DistilGPT2 | Small, fast, suitable for command interpretation and response generation. |
| TTS | pyttsx3 | Lightweight, offline TTS; integrates easily with Python. |
| Hardware | Python simulation | HardwareController simulates GPIO actions; can extend to Raspberry Pi / Arduino. |
| Audio | Sound device, simple audio | Used for microphone capture and playback in local environment. |

3) Constraints & Trade-offs

| Constraint | Discussion |
|-------------------|--|
| Model Size | Whisper-small (~74 MB) and DistilGPT2 (~332 MB) fit local CPU with limited RAM. |
| Latency | ASR + SLM + TTS response ~1–3 seconds; suitable for simple home commands. |
| RAM Usage | Models run on 8–16 GB CPU machines; small models prevent memory overload. |
| Power Needs | Entire system can run on a laptop; optional Raspberry Pi/Arduino can extend to low-power hardware. |
| Accuracy vs Speed | Whisper-small is accurate for simple commands; DistilGPT2 may require careful prompt design to avoid repetitive outputs. |
| Offline Mode | All components work fully offline after downloading models to cache. |

4) Example Use Cases

(a) Voice-Controlled Light / Fan Simulation

- Commands like “Turn on light” or “Turn off fan” are processed.
- HardwareController simulates the action with logs and optional GPIO control.
- Response is played via TTS, e.g., “Turning on the light.”

(b) Local Command-Line Voice Assistant

- Runs entirely offline.
- Users speak commands; assistant responds via audio and logs actions.

- Demonstrates end-to-end AI pipeline: ASR → SLM → TTS → Action.

(c) Smart Home Prototype

- Extendable to Raspberry Pi GPIO for controlling lights, fans, or other devices.
- Audio logs maintained for debugging and analytics.

5) Notes & Recommendations

- **Model Cache:** Pre-download Whisper & DistilGPT2 into ~/.cache/huggingface for offline use.
- **Logging:** Logger class stores speech, hardware actions, and errors for traceability.
- **Prompt Design:** Keep SLM prompts structured to avoid repetition.
- **Hardware Extension:** Later add real GPIO via Raspberry Pi or Arduino using pySerial.
- **TTS Quality:** pyttsx3 is lightweight but lower quality; Coqui/Silero can replace if higher fidelity is needed.